

# *Quora Insincere Questions Classification*

Springboard Data Science Career Track Program  
Capstone Project # 2

Author : Ashish Mohan Sharma  
Reviewer: Kenneth  
Publish Data: 07/07/2019

## Table of Contents

<b>Introduction:</b>	<b>3</b>
<b>Problem Statement:</b>	<b>3</b>
<b>Potential Clients:</b>	<b>4</b>
<b>Data Acquisition and Understanding:</b>	<b>4</b>
<b>Metrics:</b>	<b>5</b>
<b>Exploratory Data Analysis:</b>	<b>6</b>
Data Exploration:	6
Let's explore some of the Sincere and Insincere Questions and try to infer the underlying pattern of both classes intuitively.	7
Class Imbalance:	8
Most Common First Word	9
Question Length Distribution:	10
Outlier Detection:	10
Word-Cloud Visualization:	12
StopWords Distribution	13
Bigram/Trigram Plots:	13
<b>Pre-Processing:</b>	<b>17</b>
<b>WordCloud After Data Cleaning:</b>	<b>20</b>

## **Introduction:**

[Quora](#) is a platform that empowers people to learn from each other. On Quora, people can ask questions and connect with others who contribute unique insights and quality answers. Sincere questions are inquiries about which individuals genuinely want to know an answer or gain information from rather than argue a point or make a statement. For instance, “What are some sleep hacks and tips?” is deemed as sincere since it is a candid question which is most likely asked by an individual who would like some advice on falling asleep. On the other hand, insincere questions intend to make some sort of a statement and are usually asked not with the intention of receiving a helpful answer/comment. It is often the case that insincere questions target religion, gender, politics, etc. and are constructed in a non-neutral tone, are exaggerated, or use words that attack various groups. For example, "If blacks support school choice and mandatory sentencing for criminals why don't they vote Republican?" is classified as insincere since it targets the Black community and is phrased in a demeaning manner. A key challenge is to weed out insincere questions -- those founded upon false premises, or that intend to make a statement rather than look for helpful answers.

Here I will try to develop models that identify and flag insincere questions. To date, Quora has employed both machine learning and manual review to address this problem. I will try to help Quora to develop more scalable methods to detect toxic and misleading content and try to combat online trolls at scale. I will try to help Quora uphold their policy of “Be Nice, Be Respectful” and continue to be a place for sharing and growing the world’s knowledge.

## **Problem Statement:**

The goal is to handle unethical comments, in order to improve online conversation on Quora. This is done by identifying them and flagging them as insincere question. An insincere question is defined as a question intended to make a statement rather than look for helpful answers.

The process for solving such problem would be to use a supervised machine learning algorithm (many of them will be tested in this project) that will learn the patterns from

the labelled comments. The final model is supposed to, given a random comment, be able to classify it as ethical (good) or unethical (insincere).

## **Potential Clients:**

Though our model will be specific to Quora but this specific work can be utilized in other Social Media platforms as well to make sure the comments are Ethical and Sincere to be posted on the platform. The model can be integrated as first filter level even before submitting the comment. Those questions that are flagged as insincere wouldn't be allowed to be submitted at all.

It will help Social Media Platforms like Facebook, Instagram, WhatsApp etc to make sure that their users are not bullied or terrified over their system by anyone. That will also help them in creating their own credibility.

## **Data Acquisition and Understanding:**

The data I have used comes from a Kaggle competition and is separated into two datasets, training and testing. The training set consists of a list of questions along with their target values (0 for sincere and 1 for insincere). However, the testing data only contains the sample questions without the target values.

Data can be downloaded from the below link given below.

<https://www.kaggle.com/c/quora-insincere-questions-classification/data>

An insincere question is defined as a question intended to make a statement rather than look for helpful answers. Some characteristics that can signify that a question is insincere:

- Has a non-neutral tone
  - Has an exaggerated tone to underscore a point about a group of people
  - Is rhetorical and meant to imply a statement about a group of people
- Is disparaging or inflammatory
  - Suggests a discriminatory idea against a protected class of people, or seeks confirmation of a stereotype



- Makes disparaging attacks/insults against a specific person or group of people
  - Based on an outlandish premise about a group of people
  - Disparages against a characteristic that is not fixable and not measurable
- Isn't grounded in reality
  - Based on false information, or contains absurd assumptions
- Uses sexual content (incest, bestiality, pedophilia) for shock value, and not to seek genuine answers

In addition to this training and testing data, word embedding files (.txt) are used in the final model of this project. This file represents a collection of words with their corresponding embed vector, i.e., the number of features that represent similarities between words. Note: These vectors were obtained by training a neural network (using onehot encoded vectors as inputs), and taking the weights between the input and first hidden layer of the network after training. The embedding matrix here has a size of 2,196,017 by 300. This means that it contains 2,196,017 words, each having a vector of size 300.

I am interested in this specific dataset and the classification of questions as sincere or insincere because we are all active users of Quora. Additionally, I understand that reducing unproductive and harmful content online is a challenge for many tech companies including Quora. This motivated me to gain perspective on the process of moderating online speech.

## **Metrics:**

The F1 score is used for this project to evaluate the performance of the models. While accuracy is a good metric, it will not show a real indication of how the model is performing. Especially since the labels, as we will see in the next section, are highly imbalanced. The model could be saying everything is insincere and still get a high accuracy. F1 score considers both the precision and the recall of the test to compute the score: precision is the number of correct positive results divided by the number of all positive results returned by the model, and recall is the number of correct positive results

divided by the number of all samples that should have been identified as positive. The F1 score is the harmonic average of the precision and recall, where the best score is 1 and worst is 0.

**Formula:**  $2*((\text{precision}*\text{recall})/(\text{precision}+\text{recall}))$ .

## **Exploratory Data Analysis:**

It is a good practice to understand the data first and try to gather as many insights from it. EDA is all about making sense of data in hand, before getting them dirty with it.

Exploratory Data Analysis (EDA) is an approach/philosophy for data analysis that employs a variety of techniques (mostly graphical) to

1. maximize insight into a data set;
2. uncover underlying structure;
3. extract important variables;
4. detect outliers and anomalies;
5. test underlying assumptions;
6. develop parsimonious models; and
7. determine optimal factor settings.

It is an approach to data analysis that postpones the usual assumptions about what kind of model the data follow with the more direct approach of allowing the data itself to reveal its underlying structure and model.

## **Data Exploration:**

As I explained earlier, we have 2 datasets of which one is Training data while other is Test data. Train data is labeled whether the question is Sincere or Insincere while Test data is not labeled.

### **Sample of Training Dataset**

qid	question_text	target
00002165364db923c7e6	How did Quebec nationalists see their province...	0

## Q Quora Insincere Question Classification Q

000032939017120e6e44	Do you have an adopted dog, how would you enco...	0
0000412ca6e4628ce2cf	Why does velocity affect time? Does velocity a...	0
000042bf85aa498cd78e	How did Otto von Guericke used the Magdeburg h...	0
0000455dfa3e01eae3af	Can I convert montra helicon D to a mountain b...	0

Looking at the training data I found that there is not data missing in the whole dataset. Also noticed that it has **1225312** records of Sincere data while **80810** of Insincere data. Below are the details of it.

<b>Number of Sincere Questions</b>	<b>1225312</b>
<b>Number of InSincere Questions</b>	<b>80810</b>
<b>% of Sincere Questions in Train Dataset</b>	<b>0.938</b>
<b>% of InSincere Questions in Train Dataset</b>	<b>0.0619</b>

Let's explore some of the Sincere and Insincere Questions and try to infer the underlying pattern of both classes intuitively.

### Sincere Questions:

1. How did Quebec nationalists see their province as a nation in the 1960s?
2. Do you have an adopted dog, how would you encourage people to adopt and not shop?
3. Why does velocity affect time? Does velocity affect space geometry?
4. How did Otto von Guericke used the Magdeburg hemispheres?
5. Can I convert montra helicon D to a mountain bike by just changing the tyres?

### Insincere Questions:

1. Has the United States become the largest dictatorship in the world?
2. Which babies are more sweeter to their parents? Dark skin babies or light skin babies?
3. If blacks support school choice and mandatory sentencing for criminals why don't they vote Republican?
4. I am gay boy and I love my cousin (boy). He is sexy, but I dont know what to do. He is hot, and I want to see his di\*\*. What should I do?
5. Which races have the smallest penis?


## Q Quora Insincere Question Classification Q

From the above insincere questions analysis, we can infer that :

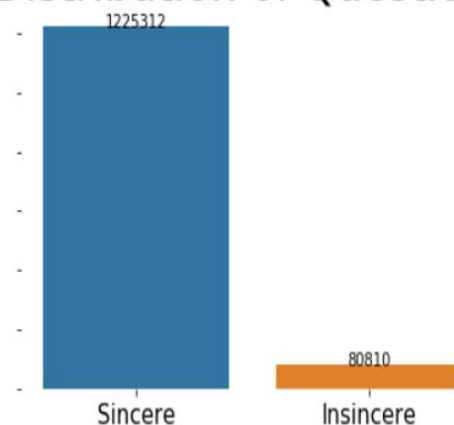
- **They have a non-neutral tone**
  - They have an exaggerated tone to underscore a point about a group of people
  - They are rhetorical and meant to imply a statement about a group of people
- **They are disparaging or inflammatory**
  - They suggests a discriminatory idea against a protected class of people, or seeks confirmation of a stereotype
  - They make disparaging attacks/insults against a specific person or group of people
  - They are based on an outlandish premise about a group of people
  - They make disparages against a characteristic that is not fixable and not measurable
- **They aren't grounded in reality.**
- **They are based on false information, or contains absurd assumptions.**
- **They use sexual content (incest, bestiality, pedophilia) for shock value, and not to seek genuine answers.**

### Class Imbalance:

% of Sincere & Insincere Questions



Distribution of Questions

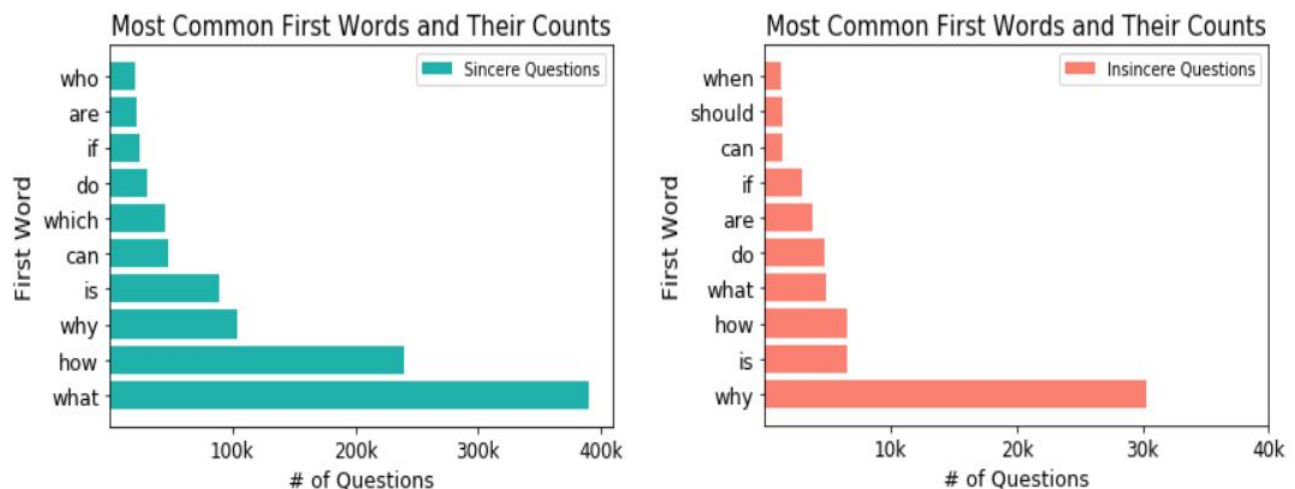




Imbalanced classes are a common problem in machine learning classification where there is a disproportionate ratio of observations in each class. With just **6.6%** of our dataset belonging to the target class, we can definitely have an imbalanced class!


This is a problem because many machine learning models are designed to maximize overall accuracy, which especially with imbalanced classes may not be the best metric to use. Classification accuracy is defined as the number of correct predictions divided by total predictions times 100. For example, if we simply predicted that all questions are sincere, we would get a classification accuracy score of **93%**!

## Most Common First Word



There seem to be not much of a difference between the Sincere and Insincere Questions if we analyse most common first words . 10 most frequent words are almost similar like *What, How, Are, Why*. This is intuitive as well, as all are questions so these words should be the part of this result.

## Question Length Distribution:



Insincere questions seem to be longer on average and have a larger variance on their length. Sincere questions seem to clump up around word length as 9 while insincere have a less pronounced peak at around 10 words.

When an individual is hoping to assert an opinion, it generally takes more words to do this than asking a question. We noticed that among insincere questions, users were often attempting to state an opinion or make discriminatory remarks about a group. Perhaps, this is why we see the above observation.

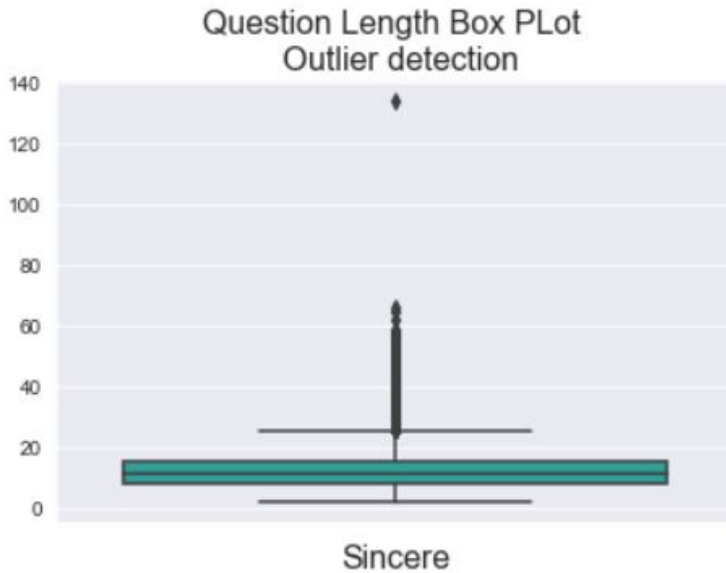
## Outlier Detection:

There is an outlier for question length in the Sincere Question group. This question is 134 words long and this user asked 7 questions in one post which resulted in the long length.

Longest Sincere Question is: 134 words

*in "star trek 2013" why did they : \*spoilers\* \*spoilers\* \*spoilers\* \*spoilers\* 1)make warping look quite a bit like an hyperspace jump 2)what in the world were those bright particles as soon as they jumped. 3)why in the world did they make it possible for two entities to react in warp space in separate jumps. 4)why did spock get emotions for this movie. 5)what was the point of hiding the "enterprise" underwater. 6)when they were intercepted by the dark ship, how come they reached earth when they were far away from her. (i don't seem to remember the scene where they warp to earth). 7)how did the ship enter earth's atmosphere when it wasn't even in orbit. 8)when scotty opened the door of the black ship , how come pike and khan didn't slow down?*

Sincere Questions length summary and Box Plot:




count	1225312
mean	12.5085
std	6.75069
min	2
25%	8
50%	11
75%	15
max	134

Outliers do not seem significant in insincere questions. The longest question length for this type is roughly half as long as the longest sincere question outlier.

Longest Sincere Question is: 64 words

*to you, does being a christian mean inviting in the spirit of jesus into you and suppressing our own spirit? 'thy will not mine' and all that? do you like living as a zombie of someone else's spirit - however perfect it may be? don't you want to experience and improve your own will and spirit and live your life as you, not jesus?*

InSincere Questions length summary and Box Plot:




count	80810
mean	17.2778
std	6.75069
min	1
25%	10
50%	15
75%	23
max	64

### Word-Cloud Visualization:


Word-Cloud Visualization highlights important textual data points, it can make dull data sizzle and immediately convey crucial information.

Sincere Word-Cloud:




*Best*, *One*, *Good*, *Think* are really positive words and therefore are categorized as Sincere Questions

### Insincere Word-Cloud:



*People, Muslim, Trump , Women, girl, hate, sex, racist, gay, lie* are the major constituents of the Insincere questions.

## StopWords Distribution



Number of stopwords present in the sentence can be a great feature for classifying the Question as Insincere. The above distribution suggests that in Insincere questions, there are a lot more stopwords present. It might be due to more explanation about the negative point they want to convey.


## Bigram/Trigram Plots:

The plots for top bigrams/trigrams in insincere questions can help us determine what kind of questions Quora would like to limit or ban. People are always finding creative new ways to ask insincere questions, and it is important for Quora to be aware of any new forms of insincere questions so as to make changes to its policies and guidelines if need be. The top bigram/trigram plots are one way to group insincere questions in a more intelligible format by pinpointing frequent topics of insincere questions.

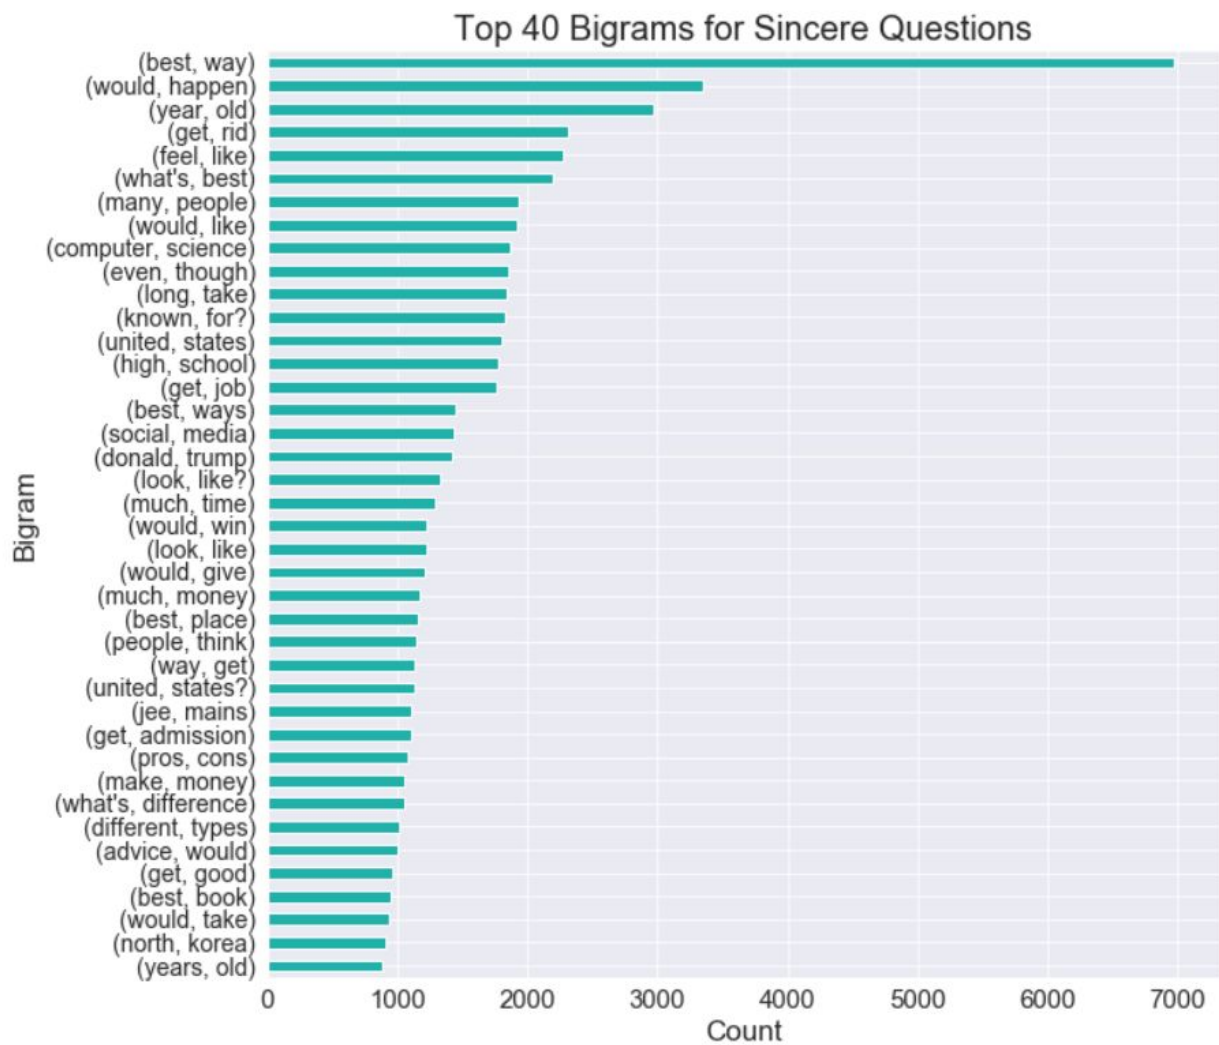
There are many potential uses for such findings including this non-exhaustive list:





- They can be examined repeatedly to regularly refine Quora's definition of an insincere question.
- They can give insight into the level of racism, sexism, and other discriminatory thoughts that are prevalent in the world. Note that although the group of users on Quora is a sample of the population of the world, the approximation for population parameters is not necessarily holistic since the set of all Quora questions is a nonrandom sample. Nonetheless, if our goal is to limit all forms of existing discriminatory thoughts, then using Quora for such an agenda is a step in our intended direction.
- Given that questions on Quora can be written anonymously, a grouping of the most frequent topics in anonymously proposed insincere questions can illustrate what some of our deepest and/or darkest thoughts are. This demonstrates that anonymity is a key factor to confession since we are often hesitant towards admitting dishonorable thoughts.



Donald Trump, Black people, white people,Indian Muslims are among the most targeted people/Community for insincere comments







## Pre-Processing:

Above analysis of data suggests that Questions contain contraction and should be cleaned. Also it has many references to other languages other than English, so need to remove those words for better results. The pre trained stopword dictionary is not the best approach as well as all stopwords not be removed. It may harm the actual meaning of the sentence.

Text can come in a variety of forms from a list of individual words, to sentences to multiple paragraphs with special characters (like tweets for example). Like any data science problem, understand the questions that are being asked will inform what steps may be employed to transform words into numerical features that work with machine learning algorithms.

Data preprocessing consists of a number of steps, any number of which may or may not apply to a given task, but generally fall under the broad categories of:

- **Tokenization** : Tokenization is a step which splits longer strings of text into smaller pieces, or tokens. Larger chunks of text can be tokenized into sentences, sentences can be tokenized into words, etc. Further processing is generally performed after a piece of text has been appropriately tokenized. Tokenization is also referred to as text segmentation or lexical analysis. Sometimes segmentation is used to refer to the breakdown of a large chunk of text into pieces larger than words (e.g. paragraphs or sentences), while tokenization is reserved for the breakdown process which results exclusively in words.

- **Normalization:**

Before further processing, text needs to be normalized. Normalization generally refers to a series of related tasks meant to put all text on a level playing field: converting all text to the same case (upper or lower), removing punctuation,

converting numbers to their word equivalents, and so on. Normalization puts all words on equal footing, and allows processing to proceed uniformly.

- **Replace contractions :**

This replace all contractions with the full meaning. (Contractions scrapped from Wikipedia:

[https://en.wikipedia.org/wiki/Wikipedia:List\\_of\\_English\\_contractions](https://en.wikipedia.org/wiki/Wikipedia:List_of_English_contractions))

Contractions are shortened version of words or syllables. They often exist in either written or spoken forms in the English language. These shortened versions or contractions of words are created by removing specific letters and sounds. In case of English contractions, they are often created by removing one of the vowels from the word. Examples would be, *do not* to *don't* and *I would* to *I'd*. Converting each contraction to its expanded, original form helps with text standardization.

- **Stemming/Lemmatization:**

Stemming and lemmatization are text normalization techniques that help reduce the number of features generated by reducing the inflectional forms of each word into a common base or root. Stemming and lemmatization differ in that stemming cuts off the end or the beginning of the word, while lemmatization takes into consideration the morphological analysis of the words. For example, 'studies' and 'studying' would stem to 'studi' and 'study' (respectively), but would lemmatize to 'study'

- **Lowering Case:**

This avoids having multiple copies of the same words. For example, while calculating the word count, ‘Analytics’ and ‘analytics’ will be taken as different words.

- **Remove Numbers:**

Replacing all number with the string ‘#’. This step was done because the total vocabulary of numbers alone was too large and most of them are non-meaningful.

- **Remove Punctuation:**

Punctuation doesn’t add any extra information while treating text data. Therefore removing all instances of it will help us reduce the size of the training data.

- **Spelling Correction:**

Spelling correction is a useful pre-processing step because this will also help us in reducing multiple copies of words. For example, “*Analytics*” and “*analytcs*” will be treated as different words even if they are used in the same sense. Common misspellings scrapped from oxford dictionaries:

<https://en.oxforddictionaries.com/spelling/common-misspellings>)

- **Strip white space**

- **Remove Stop Words:**

Stop words (or commonly occurring words) should be removed from the text data. For this purpose, we can either create a list of stopwords ourselves or we can use predefined libraries. We have to be careful while removing the Stopwords as they can change the meaning of the sentence as well.

- **Remove 25 commonly occurring words and 25 most rare words from corpus**

## World Cloud of Sincere Questions:

